

# Star Space

September 24, 2021

## 0.1 Adventure with StarSpace: A Neural Embedding Approach

-> StarSpace is a general-purpose neural model for efficient learning of entity embeddings for

- > Learning word, sentence or document level embeddings.
- > Information retrieval: ranking of sets of entities/documents or objects, e.g. ranking w
- > Text classification, or any other labeling task.
- > Metric/similarity learning, e.g. learning sentence or document similarity.
- > Content-based or Collaborative filtering-based Recommendation, e.g. recommending music o
- > Embedding graphs, e.g. multi-relational graphs such as Freebase.
- > Image classification, ranking or retrieval (e.g. by using existing ResNet features).

## 0.2 Advanced solution: StarSpace embeddings

Now you are ready to train your own word embeddings! In particular, you need to train embeddings specially for our task of duplicates detection. Unfortunately, StarSpace cannot be run on Windows and we recommend to use provided [docker container](#) or other alternatives. Don't delete results of this task because you will need it in the final project.

### 0.2.1 How it works and what's the main difference with word2vec?

The main point in this section is that StarSpace can be trained specifically for some tasks. In contrast to word2vec model, which tries to train similar embeddings for words in similar contexts, StarSpace uses embeddings for the whole sentence (just as a sum of embeddings of words and phrases). Despite the fact that in both cases we get word embeddings as a result of the training, StarSpace embeddings are trained using some supervised data, e.g. a set of similar sentence pairs, and thus they can better suit the task.

```
[36]: import pandas as pd #Analysis
import matplotlib.pyplot as plt #Visulization
import seaborn as sns #Visulization
import numpy as np #Analysis
from scipy.stats import norm #Analysis
from sklearn.preprocessing import StandardScaler #Analysis
from scipy import stats #Analysis
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
import gc
```

```

import os
import string
color = sns.color_palette()

%matplotlib inline

from plotly import tools
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.express as px

```

```

[150]: df_train = pd.read_csv('data/drugsComTrain_raw.csv')
df_test = pd.read_csv('data/drugsComTest_raw.csv')

```

```

[151]: df_train.head()

```

```

[151]:  uniqueID      drugName      condition \
0      206461      Valsartan  Left Ventricular Dysfunction
1       95260      Guanfacine                ADHD
2       92703        Lybrel          Birth Control
3      138000      Ortho Evra          Birth Control
4       35696  Buprenorphine / naloxone  Opiate Dependence

                                review  rating      date \
0  "It has no side effect, I take it in combinati...      9  20-May-12
1  "My son is halfway through his fourth week of ...      8  27-Apr-10
2  "I used to take another oral contraceptive, wh...      5  14-Dec-09
3  "This is my first time using any form of birth...      8   3-Nov-15
4  "Suboxone has completely turned my life around...      9  27-Nov-16

    usefulCount
0              27
1             192
2              17
3              10
4              37

```

```

[152]: df_test.head()

```

```

[152]:  uniqueID      drugName      condition \
0      163740      Mirtazapine      Depression
1      206473      Mesalamine  Crohn's Disease, Maintenance
2      159672      Bactrim      Urinary Tract Infection
3       39293      Contrave      Weight Loss
4       97768  Cyclofam 1 / 35      Birth Control

```

	review	rating	date	\
0	"I've tried a few antidepressants over th...	10	28-Feb-12	
1	"My son has Crohn's disease and has done ...	8	17-May-09	
2	"Quick reduction of symptoms"	9	29-Sep-17	
3	"Contrave combines drugs that were used for al...	9	5-Mar-17	
4	"I have been on this birth control for one cyc...	9	22-Oct-15	


	usefulCount
0	22
1	17
2	3
3	35
4	4

```
[153]: df_all = pd.concat([df_train,df_test]).reset_index()
del df_all['index']
```

```
[154]: df_all = df_all.dropna(axis=0)
df_all = df_all.drop_duplicates(subset=['uniqueID','review'])
```

```
[155]: condition_dn = df_all.groupby(['condition'])['drugName'].nunique().
↳sort_values(ascending=False)
condition_dn[0:20].plot(kind="bar", figsize = (14,6), fontsize = 10, color="green")
plt.xlabel("", fontsize = 20)
plt.ylabel("", fontsize = 20)
plt.title("Top20 : The number of drugs per condition.", fontsize = 20)
```

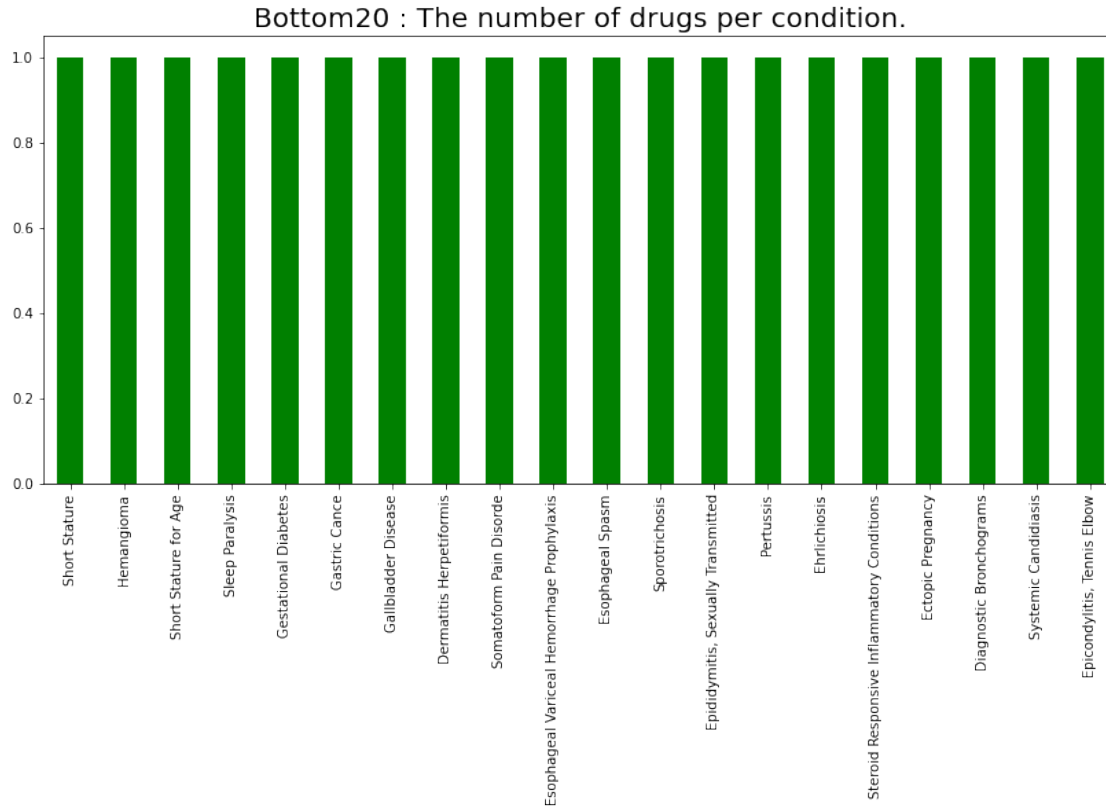
```
[155]: Text(0.5, 1.0, 'Top20 : The number of drugs per condition.')
```



```
[156]: condition_dn = df_all.groupby(['condition'])['drugName'].nunique().
        ↳sort_values(ascending=False)

condition_dn[condition_dn.shape[0]-20:condition_dn.shape[0]].plot(kind="bar",
        ↳figsize = (14,6), fontsize = 10,color="green")
plt.xlabel("", fontsize = 20)
plt.ylabel("", fontsize = 20)
plt.title("Bottom20 : The number of drugs per condition.", fontsize = 20)
```

```
[156]: Text(0.5, 1.0, 'Bottom20 : The number of drugs per condition.')
```



```
[157]: count_df = df_all[['condition', 'review']].groupby('condition').
        ↳aggregate({'review': 'count'}).reset_index().
        ↳sort_values('review', ascending=False)
count_df.head()
```

```
[157]:
```

	condition	review
175	Birth Control	38436
273	Depression	12164
613	Pain	8245
133	Anxiety	7812
87	Acne	7435

```
[158]: target_conditions = count_df[count_df['review'] > 3000]['condition'].values
```

```
[159]: def condition_parser(x):
        if x in target_conditions:
            return x
        else:
            return "OTHER"

df_all['condition'] = df_all['condition'].apply(lambda x: condition_parser(x))
```

```
[160]: df_all = df_all[df_all['condition'] != 'OTHER']
```

```
[161]: px.bar(count_df[count_df['review'] > 3000], x='condition', y='review')
```

```
[162]: import contractions
```

```
[163]: import re
```

```
def clean_text(x):
    pattern = r'[~a-zA-z0-9\s]'
    text = re.sub(pattern, ' ', x)
    return x

def clean_numbers(x):
    if bool(re.search(r'\d', x)):
        x = re.sub('[0-9]{5}', '#####', x)
        x = re.sub('[0-9]{4}', '####', x)
        x = re.sub('[0-9]{3}', '###', x)
        x = re.sub('[0-9]{2}', '##', x)
    return x
```

```
[164]: # lower the text
df_all["review"] = df_all["review"].apply(lambda x: x.lower())

# Clean the text
df_all["review"] = df_all["review"].apply(lambda x: clean_text(x))

# Clean numbers
df_all["review"] = df_all["review"].apply(lambda x: clean_numbers(x))

# Clean Contractions
df_all["review"] = df_all["review"].apply(lambda x: " ".join([contractions.
    ↪fix(word) for word in x.split()])))
```

```
[165]: df_all.head()
```

```
[165]:
```

	uniqueID	drugName	condition	\
1	95260	Guanfacine	ADHD	
2	92703	Lybrel	Birth Control	
3	138000	Ortho Evra	Birth Control	
6	165907	Levonorgestrel	Emergency Contraception	
7	102654	Aripiprazole	Bipolar Disorde	

	review	rating	date	\
1	"my son is halfway through his fourth week of ...	8	27-Apr-10	
2	"i used to take another oral contraceptive, wh...	5	14-Dec-09	
3	"this is my first time using any form of birth...	8	3-Nov-15	

6	"he pulled out, but he cummed a bit in me. i t...	1	7-Mar-17
7	"abilify changed my life. there is hope. i was...	10	14-Mar-15

	usefulCount
1	192
2	17
3	10
6	5
7	32

```
[196]: f = open('data.txt', 'a')
      for row in df_all.itertuples():
          f.write(row.review)
          f.write("\n")
```

```
[197]: import sentencepiece as spm
```

```
[198]: spm.SentencePieceTrainer.train(input='data.txt', model_prefix='m',
                                     vocab_size=10000)
```

```
[199]: sp = spm.SentencePieceProcessor(model_file='m.model')
```

### 0.3 Classification: Can you predict the patient's condition based on the review?

```
[200]: df_all.condition.unique()
```

```
[200]: array(['ADHD', 'Birth Control', 'Emergency Contraception',
             'Bipolar Disorder', 'Depression', 'Obesity', 'Insomnia',
             'Vaginal Yeast Infection', 'Pain', 'Diabetes, Type 2', 'Anxiety',
             'Acne', 'High Blood Pressure', 'Weight Loss'], dtype=object)
```

```
[201]: LABELS = ['ADHD', 'Birth Control', 'Emergency Contraception',
                'Bipolar Disorder', 'Depression', 'Obesity', 'Insomnia',
                'Vaginal Yeast Infection', 'Pain', 'Diabetes, Type 2', 'Anxiety',
                'Acne', 'High Blood Pressure', 'Weight Loss']

EMPTY_ID = len(LABELS)

def create_labeled_string(row):
    parts = sp.encode(row["review"], out_type = 'str')
    parts.append("__label__{}".format(row['condition']))
    return " ".join(parts)
```

```
[202]: from sklearn.model_selection import train_test_split
      train, test = \
      ↪ train_test_split(df_all[['review', 'condition']], stratify=df_all['condition'],
```

```
test_size=0.25)
```

```
[203]: train_lines = train.apply(create_labeled_string, 1)
with open("cache/train.txt", "w") as f:
    f.write("\n".join(train_lines))
val_lines = test.apply(create_labeled_string, axis=1)
with open("cache/val.txt", "w") as f:
    f.write("\n".join(val_lines))
```

```
[214]: ! /Users/subir/Starspace/./starspace train -thread 4 -trainFile cache/train.
      ↪txt -model cache/starspace.modela
```

Arguments:

```
lr: 0.01
dim: 100
epoch: 5
maxTrainTime: 8640000
validationPatience: 10
saveEveryEpoch: 0
loss: hinge
margin: 0.05
similarity: cosine
maxNegSamples: 10
negSearchLimit: 50
batchSize: 5
thread: 4
minCount: 1
minCountLabel: 1
label: __label__
label: __label__
ngrams: 1
bucket: 2000000
adagrad: 1
trainMode: 0
fileFormat: fastText
normalizeText: 0
dropoutLHS: 0
dropoutRHS: 0
useWeight: 0
weightSep: :
Start to initialize starspace model.
Build dict from input file : cache/train.txt
Read 9M words
Number of words in dictionary: 10085
Number of labels in dictionary: 14
Loading data from file : cache/train.txt
Total number of examples loaded : 83673
```



```

Training epoch 0: 0.01 0.002
Epoch: 100.0% lr: 0.008000 loss: 0.009490 eta: 0h2m tot: 0h0m41s
(20.0%)tot: 0h0m1s (0.7%) (6.8%) tot: 0h0m19s (8.8%)63.8% lr: 0.008602
loss: 0.011834 eta: 0h3m tot: 0h0m27s (12.8%)72.8% lr: 0.008506 loss:
0.011035 eta: 0h3m tot: 0h0m30s (14.6%)0h0m31s (14.8%)32s (15.1%)15.4%)s
(15.8%)94.7% lr: 0.008121 loss: 0.009712 eta: 0h2m tot: 0h0m39s (18.9%)m39s
(19.0%)
----+++ Epoch 0 Train error : 0.00939365 +++---
Training epoch 1: 0.008 0.002
Epoch: 100.0% lr: 0.006024 loss: 0.003038 eta: 0h1m tot: 0h1m21s (40.0%)8%
lr: 0.007855 loss: 0.002965 eta: 0h2m tot: 0h0m44s (21.4%)0m44s
(21.4%)0h0m45s (21.9%)49s (23.9%)33.7% lr: 0.007446 loss: 0.003067 eta:
0h2m tot: 0h0m55s (26.7%)0m55s (26.9%)50.9% lr: 0.007060 loss: 0.003032
eta: 0h2m tot: 0h1m1s (30.2%)54.1% lr: 0.006988 loss: 0.003112 eta: 0h2m
tot: 0h1m2s (30.8%)57.4% lr: 0.006964 loss: 0.003138 eta: 0h2m tot: 0h1m4s
(31.5%)s (31.9%)m21s (39.9%)
----+++ Epoch 1 Train error : 0.00302130 +++---
Training epoch 2: 0.006 0.002
Epoch: 100.0% lr: 0.004000 loss: 0.001989 eta: 0h1m tot: 0h2m0s
(60.0%)1m23s (40.9%)41.2%)m26s (42.4%)1m31s (45.1%)45.8%) (46.4%)33.3% lr:
0.005349 loss: 0.002033 eta: 0h1m tot: 0h1m34s (46.7%)s (46.7%)m38s
(48.5%)%)44s (51.7%)m tot: 0h1m45s (52.5%)72.1% lr: 0.004578 loss: 0.002022
eta: 0h1m tot: 0h1m49s (54.4%)73.1% lr: 0.004554 loss: 0.002013 eta: 0h1m
tot: 0h1m49s (54.6%)87.8% lr: 0.004241 loss: 0.002022 eta: 0h1m tot:
0h1m55s (57.6%)59s (59.7%)
----+++ Epoch 2 Train error : 0.00206770 +++---
Training epoch 3: 0.004 0.002
Epoch: 100.0% lr: 0.002000 loss: 0.001655 eta: <1min tot: 0h2m38s
(80.0%)%)25.8% lr: 0.003398 loss: 0.001564 eta: 0h1m tot: 0h2m10s
(65.2%)11s (65.9%)0h2m19s (70.0%) (70.4%)71.0% lr: 0.002386 loss: 0.001687
eta: <1min tot: 0h2m27s (74.2%)74.5%)76.2%) (77.7%) (78.7%)97.2% lr:
0.002048 loss: 0.001636 eta: <1min tot: 0h2m37s (79.4%) (79.9%)
----+++ Epoch 3 Train error : 0.00159319 +++---
Training epoch 4: 0.002 0.002
Epoch: 100.0% lr: 0.000024 loss: 0.001269 eta: <1min tot: 0h3m18s
(100.0%)54s (88.1%)4s (93.2%)76.4% lr: 0.000458 loss: 0.001296 eta: <1min
tot: 0h3m8s (95.3%)85.7% lr: 0.000217 loss: 0.001288 eta: <1min tot:
0h3m12s (97.1%) (97.6%)
----+++ Epoch 4 Train error : 0.00131867 +++---
Saving model to file : cache/starspace.model
Saving model in tsv format : cache/starspace.model.tsv

```

```

[217]: ! /Users/subir/Starspace/./starspace test -thread 4 -testFile cache/val.txt
↪-model cache/starspace.model -predictionFile cache/starspace.pred

```

Arguments:  
lr: 0.01  
dim: 100

```

epoch: 5
maxTrainTime: 8640000
validationPatience: 10
saveEveryEpoch: 0
loss: hinge
margin: 0.05
similarity: cosine
maxNegSamples: 10
negSearchLimit: 50
batchSize: 5
thread: 4
minCount: 1
minCountLabel: 1
label: __label__
label: __label__
ngrams: 1
bucket: 2000000
adagrad: 1
trainMode: 0
fileFormat: fastText
normalizeText: 0
dropoutLHS: 0
dropoutRHS: 0
useWeight: 0
weightSep: :
Start to load a trained starspace model.
STARSPACE-2018-2
Model loaded.
Loading data from file : cache/val.txt
Total number of examples loaded : 27891
Predictions use 14 known labels.
-----Loaded model args:
Arguments:
lr: 0.01
dim: 100
epoch: 5
maxTrainTime: 8640000
validationPatience: 10
saveEveryEpoch: 0
loss: hinge
margin: 0.05
similarity: cosine
maxNegSamples: 10
negSearchLimit: 50
batchSize: 5
thread: 4
minCount: 1
minCountLabel: 1

```

```
label: __label__
label: __label__
ngrams: 1
bucket: 2000000
adagrad: 1
trainMode: 0
fileFormat: fastText
normalizeText: 0
dropoutLHS: 0
dropoutRHS: 0
useWeight: 0
weightSep: :
Predictions use 14 known labels.
Evaluation Metrics :
hit@1: 0.473988 hit@10: 0.998315 hit@20: 1 hit@50: 1 mean ranks : 1.66581 Total
examples : 27891
```

[ ]: